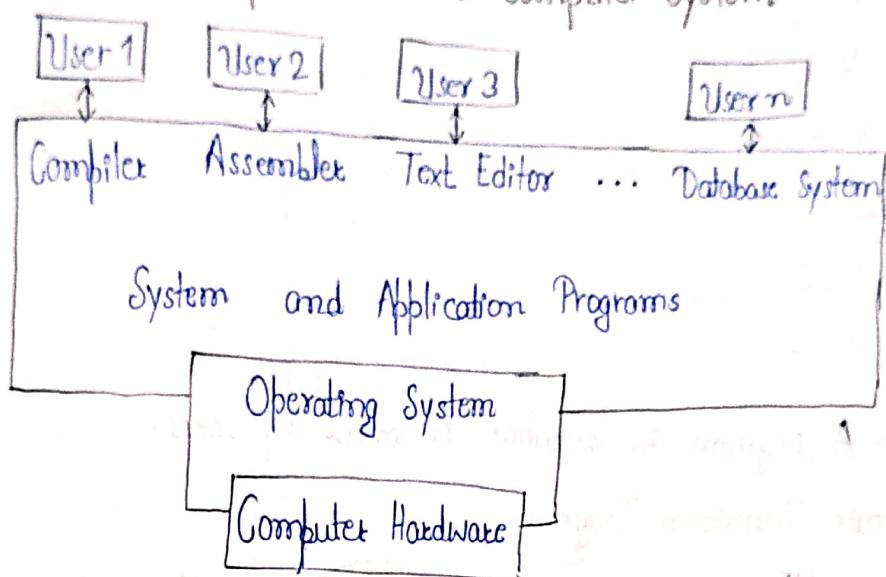
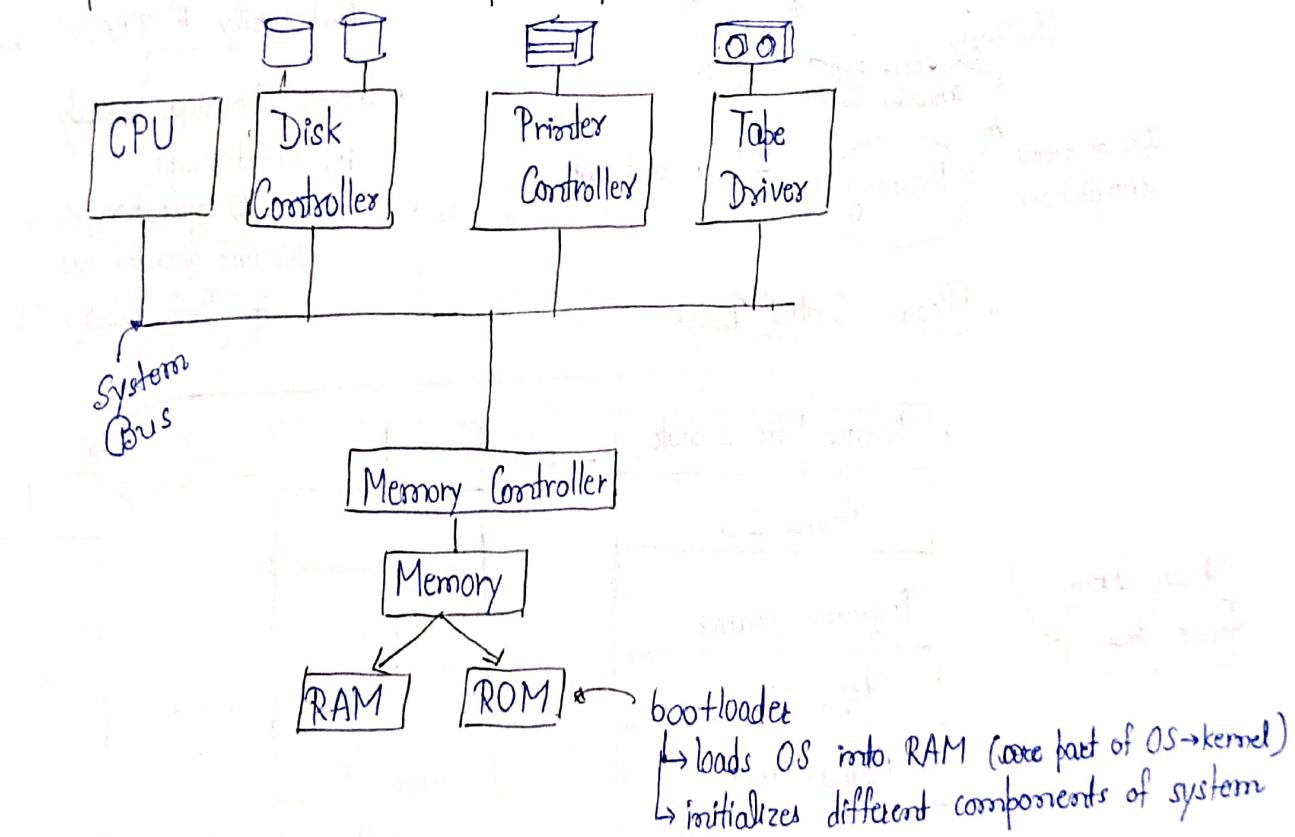


Operating System

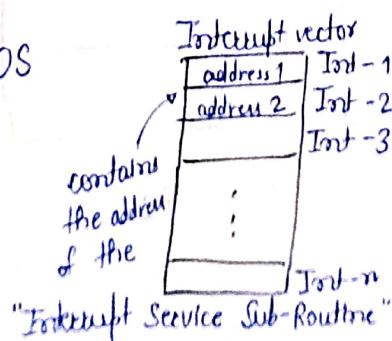
Abstract View of Components of a Computer System



* Components of Modern Computer System



* Interrupt-Driven OS



Real Time OS

OS → Batch OS → Multi-programming OS → Time-Sharing OS

- * Centralized - control comes from one program to another

Types of OS:

Batch OS

Multi-programming OS

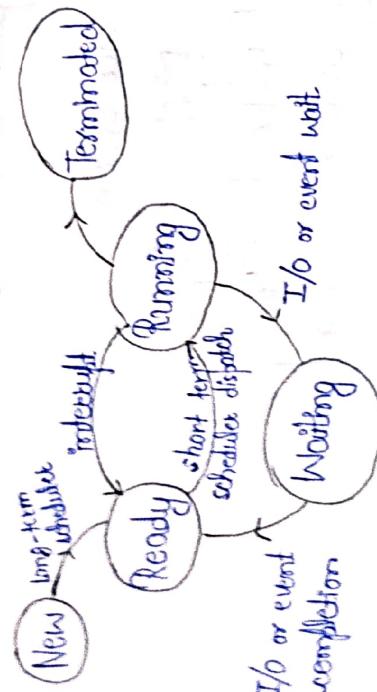
Time-Sharing OS

Real-Time OS

Distributed OS

- * Process = A program in execution is called a process.

* Process State Transition Diagram

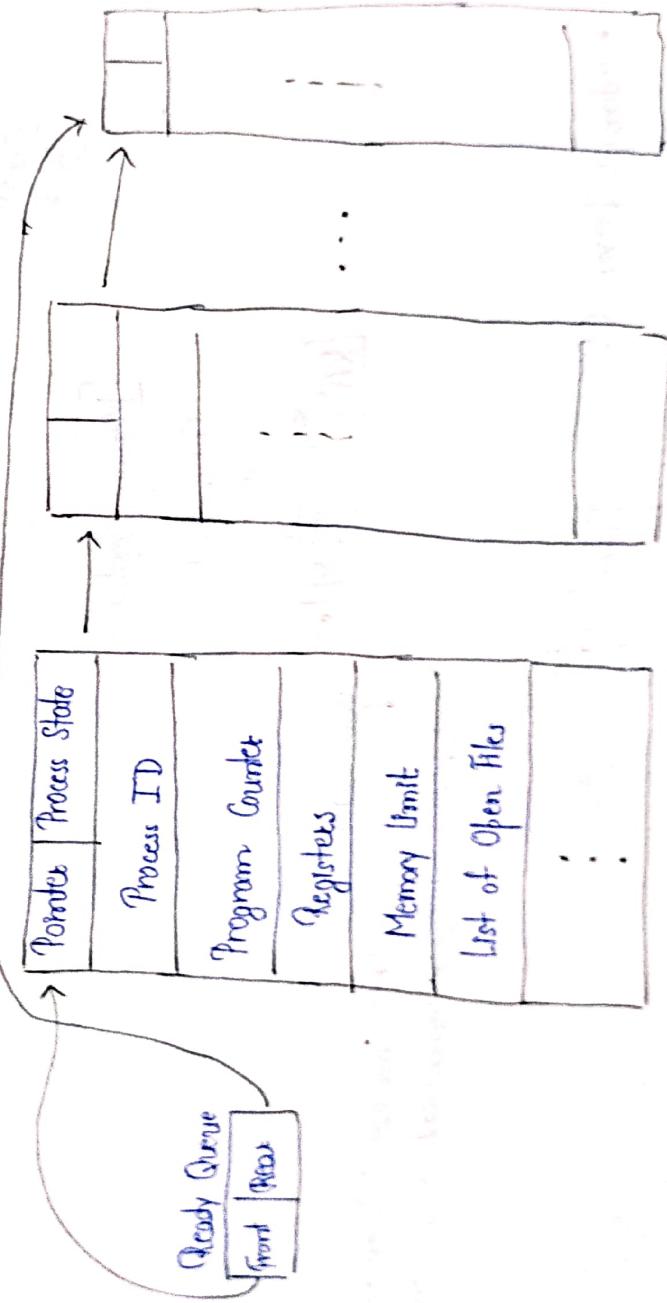


* Medium-term scheduler:
- switches out some of the processes temporarily if "READY" is excess

- * Process: Running → Ready
 - if, interrupt
 - vi) I/O operation
 - vi) time quanta over

* Process Control Block

in Time-sharing OS



P1

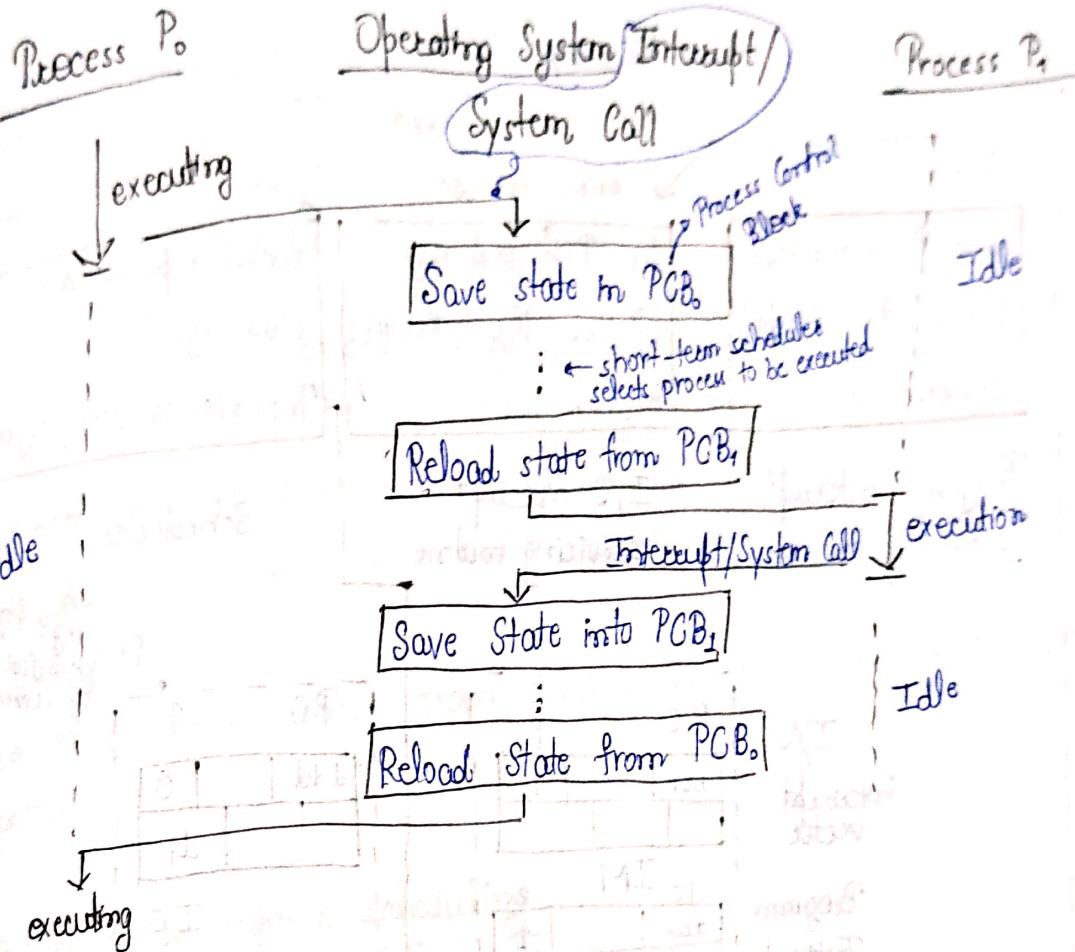
P2

Pm

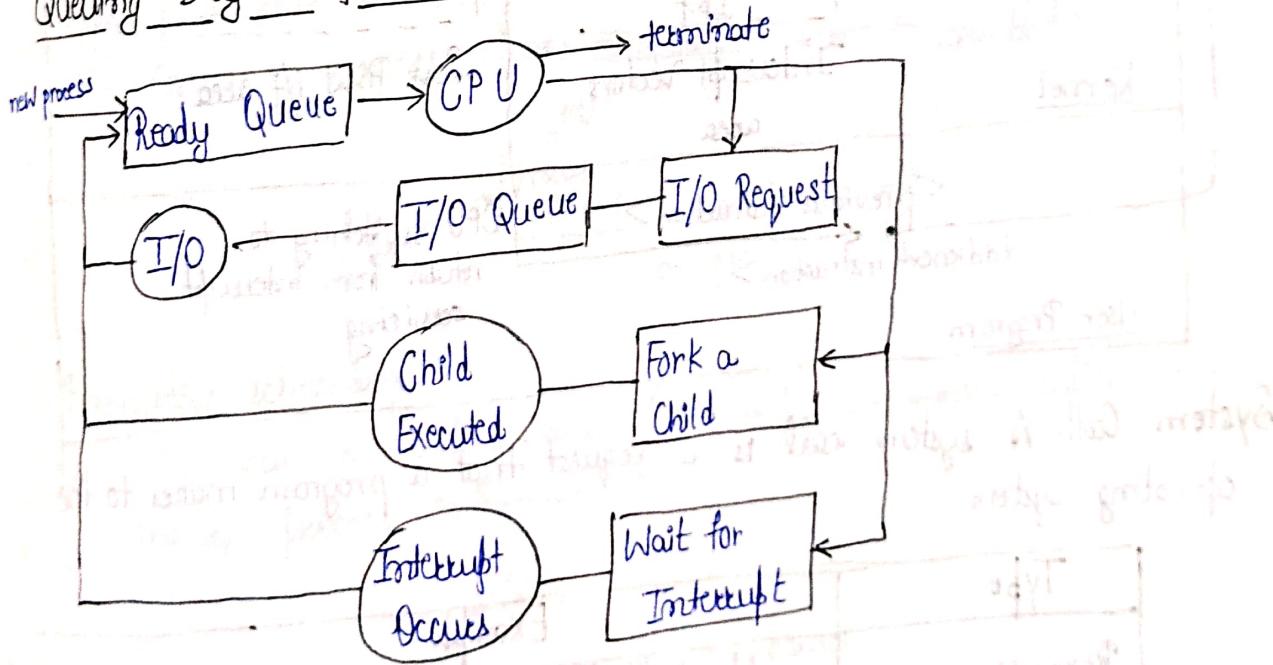
process field contains "switched"



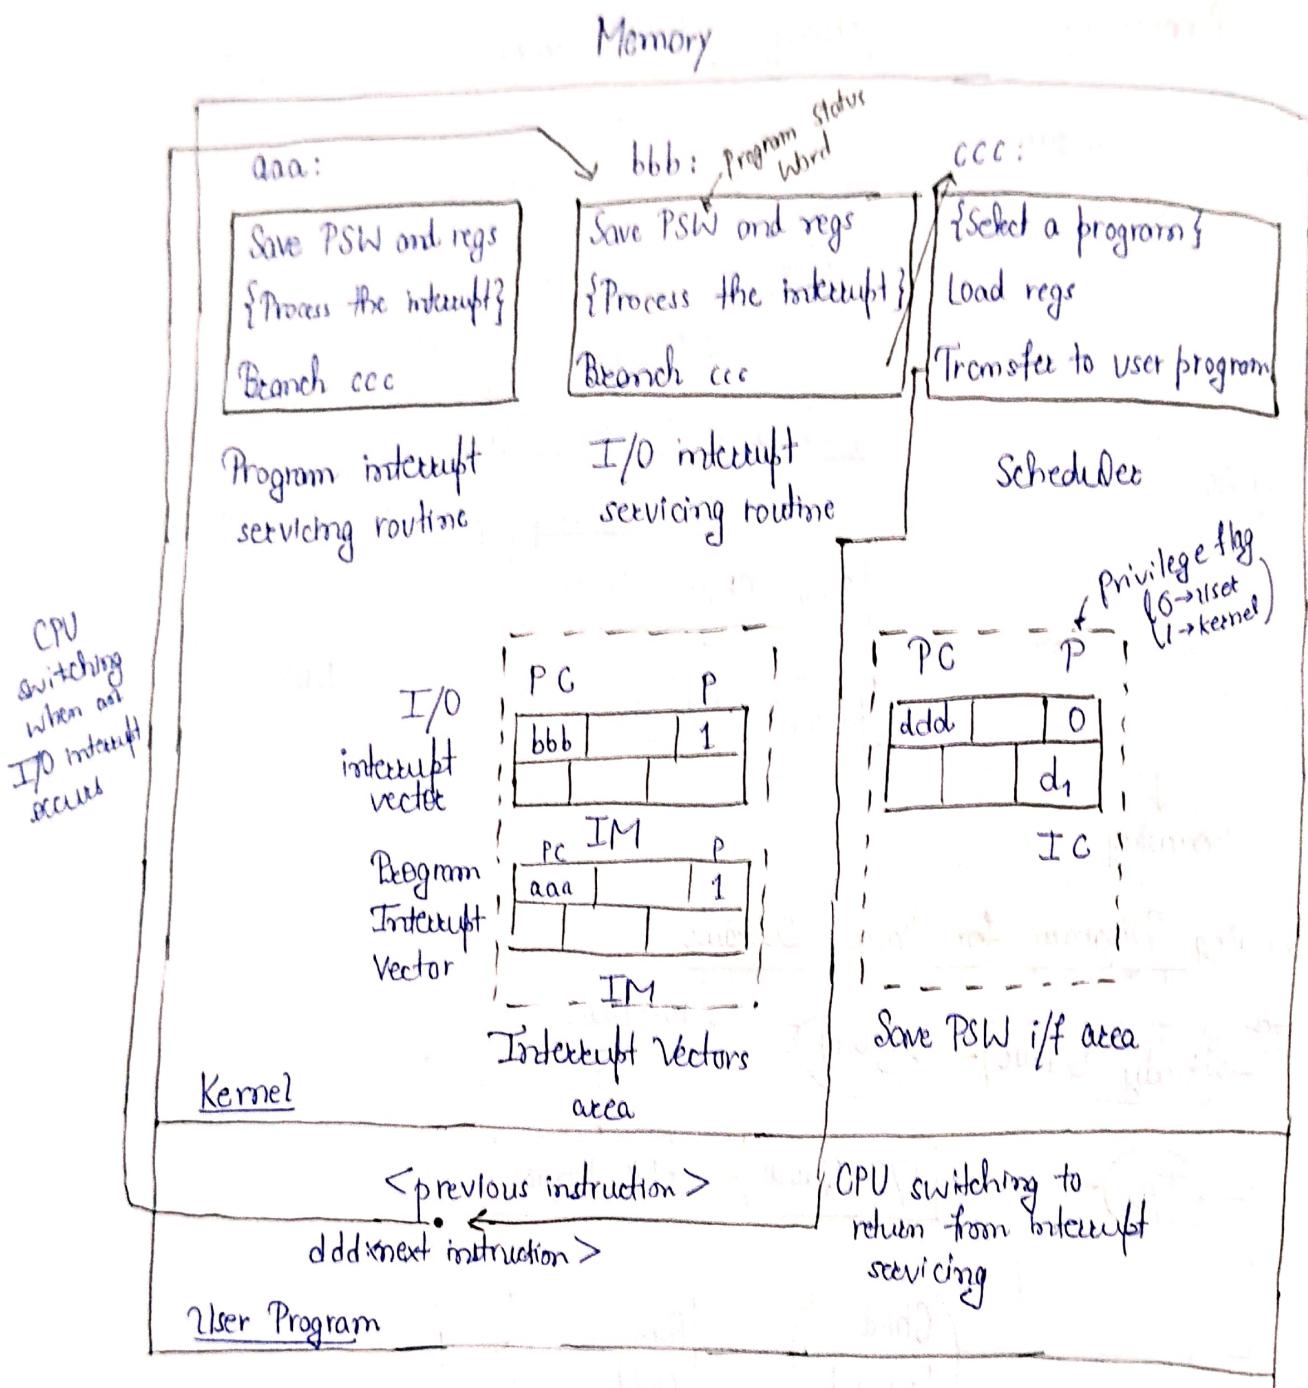
Context Switching



Queuing Diagram for Process Schedule



Interrupt Processing



System Call: A system call is a request that a program makes to the operating system.

Type	Example
Resource	allocate/deallocate resource, check resource available
Program/Process	Execute/terminate program
File	open/close file, read/write file
Information	get time information, O.S. information
Communication	send/receive message through network

Multi-Programming

- CPU-bound program
- I/O-bound program

Term Scheduler

• Controls degree of multi-programming

no. of processes running simultaneously

Throughput: no. of processes completed by CPU in unit time from submission of a process to its

Turn-around time: total time taken from submission of a process until the first completion

Response time: time taken from submission of a process until its completion

Priority-based pre-emptive scheduling

• Non-preemptive scheduling starts executing, it runs to completion

- Once a process



P₁ → [Process P₀] → P₀ won't stop until completely executed

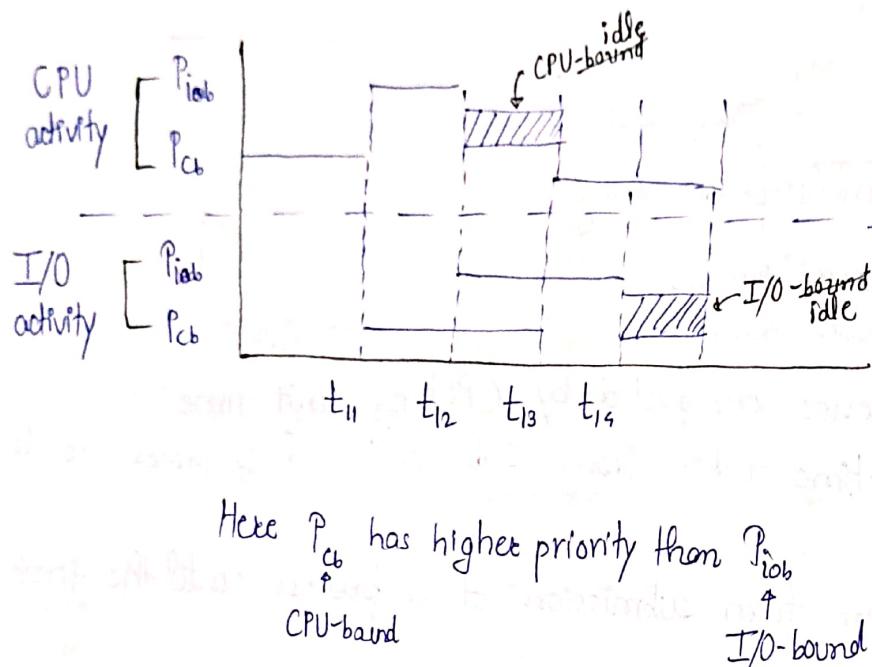
• Preemptive scheduling takes away from a process mid-execution if a higher priority process arrives

- CPU can be taken away if priority process interrupts it



P₁ → [Process P₀] → P₀ will be paused if P₁ is higher priority

Timing Chart when CPU-based bound program has higher priority



Here P_{cb} has higher priority than P_{ab}

CPU-bound

I/O-bound

Operation on Process

→ Creation of child process using fork() system call

```
#include <unistd.h>
```

```
#include <sys/types.h>
```

```
#include <stdio.h>
```

```
#include <sys/wait.h>
```

```
int main()
```

```
{ pid_t p;
```

```
printf("Before fork \n");
```

$p = \text{fork}();$ ← system call that creates a new process by duplicating the calling process

```
if ( $p == 0$ ) { //child process
```

```
printf("I am child having id %d", getpid());
```

```
printf("My parent's id is %d", getppid());
```

```
} else { //parent process
```

```
wait(null); ← parent process will wait until child process has finished
```

```
printf("My child's id is %d \n", p);
```

```
printf("I am parent having id %d \n", getpid());
```

```
}
```

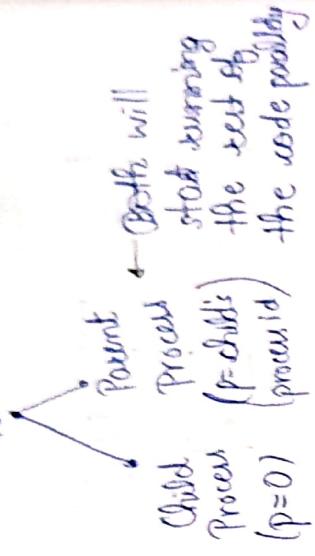
```
printf("Common \n");
```

```
} //End of main function
```

Output:

Before fork
 I am child having Id 20
 My parent's id is 15
 Common
 Child's id is 20
 My parent having Id 15
 I am common

Parent Process



- If a child process is intended to perform some subtask on a system with multiple CPUs, then execution will be sped up as compared to using a function call to perform the subtask.
- The process id is generally greater than the parent's process id (because its child's needed later)

→ Inter Process Communication using Shared Memory

#define buffer_size 10

```
typedef struct {
    - - -
}
```

```
item buffer [buffer_size];
```

```
int in=0;
```

```
int out=0;
```

```
while (1) {
```

→ Product process

```
/* Produce an item in next Produced*/
```

```
while (((in + 1) % buffer_size) == out) {
```

```
    // do nothing
```

```
    buffer [in] = nextProduced;
```

```
    in = (in + 1) % buffer_size;
```

```
}
```

```
while (1) {
```

→ Consumer process

```
while (in == out) {
```

```
    // do nothing
```

```
}
```

```
nextConsumed = buffer [out];
```

```
out = (out + 1) % buffer_size;
```

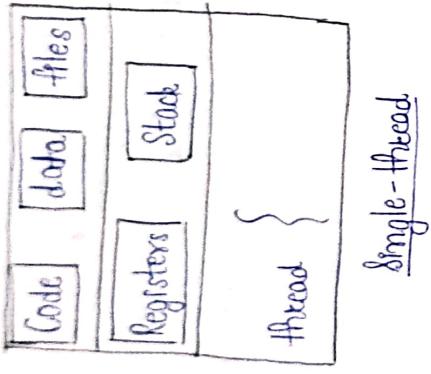
```
/* Consume the item in next consumed */
```

```
if (nextConsumed == in) {
```

```
    in = (in + 1) % buffer_size;
```

```
    /* Consume the item in next consumed */
```

Threading



Single-thread

• Thread ~ "light-weight process"

• A thread is the smallest unit of execution within a process

Benefits of Threading:

- Responsiveness
- Resource Sharing
- Economic
- Utilization of multi-processor

CPU Scheduling Techniques

- First come, first served
- Shortest Job First
- Priority
- Round-Robin

(for round-robin) Priority

Time Slice	Process	Priority
= 1	P ₁	2
	P ₂	4
	P ₃	1

Burst Time

total time a process needs the CPU for execution without any interruption

→ total time spent waiting in ready queue before it gets CPU for execution

FCFS:

	P ₁	P ₂	P ₃	
0	24	24	27	30

Avg. Waiting Time: $\frac{0+24+27}{3} = 17$ (total time a process spends waiting in ready queue before it gets CPU for execution)

Avg. Turn-around Time: $\frac{24+27+30}{3} = 27$ (total time taken from the moment a process arrives in the system to the moment it completes execution)

Gantt Chart

Diagram illustrating the Gantt chart for the processes P₁, P₂, and P₃. The chart shows the sequence of processes over time, with P₁ running for 24 units, P₂ for 24 units, P₃ for 27 units, and then returning to P₁ for the final 3 units.

Garrett Chart

	P ₂	P ₁	P ₃	P ₂	P ₁
STF:	0	3	6	12	18

Avg. Waiting Time = $\frac{2+3+6+12}{3} = 3$

Avg. Turn-around Time = $\frac{3+6+12+18}{3} = 12$

Garrett Chart

Priority:	P ₂	P ₁	P ₃	P ₂	P ₁
	0	3	6	9	12

Avg. Waiting Time = $\frac{0+3+27}{3} = 10$

Avg. Turn-around Time = $\frac{3+27+30}{3} = 20$

Gantt Chart

	P ₁	P ₂	P ₃	P ₁	P ₂	P ₃	P ₁
Round-Robin:	0	4	7	10	14	19	26

P₁ → (0+6) + 4 + 7 = 17 → $\frac{17}{3} \approx 5.67$

Avg. Waiting Time = $\frac{P_1 + P_2 + P_3}{3} = \frac{3+7+10}{3} = 6.67$

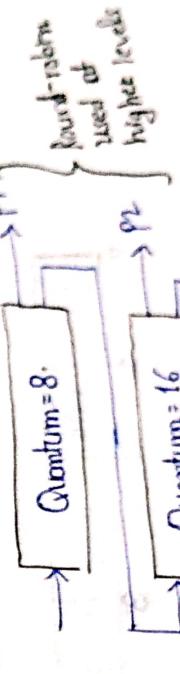
Avg.

Turn-around Time

CPU Scheduling Algorithm

(a) Multi-level Queue

(b) Multi-level Feedback Queue



P1	6
P2	12
P3	30

	Arrival Time	Priority	Burst Time
P1	0	2	4
P2	1	4	2

Non-preemptive:

6
4
0

P1	P2	P1
0	+	3

P1	P2	P1
0	+	3

- * Starvation: Lower priority processes do not get access to the CPU as higher priority processes keep coming.
 - * Aging : After waiting for some time, the process moves up to a higher priority queue.

CPU Scheduling Problem with mixed CPU burst and I/O burst

<u>Process</u>	<u>Arrival Time</u>	<u>Priority</u>	<u>CPU</u>	<u>I/O</u>	<u>CPU</u>
P1	0	2	1	5	3
P2	2	3	3	3	1
P3	3	1	2	3	1
P4	3	4	2	4	1

Model: Pre-emptive Scheduling-Priority

Draw Gantt chart.

	P1	Idle	P2	P4	P2	P1	P4	P2	P1	P3	P3	Idle	P3
0	2	2	3	5	7	9	10	11	12	11	12	14	8
1	3	3	4	6	8	10	11	12	13	12	13	15	9
2	4	4	5	7	9	11	12	13	14	13	14	16	10
3	5	5	6	8	10	12	13	14	15	14	15	17	11

Process Synchronization

- producer-consumer problem:

Item buffer[Buffer-size]; int counter=0;
int in=0, out=0;

Producer:

```
while(1) {
    counter = Buffer.size;
    while ((in+1)%Buffer.size) == out)
        ;
    buffer[in] = nextProduced;
    in = (in+1)%Buffer.size;
    counter++;
}
```

//Entry Section

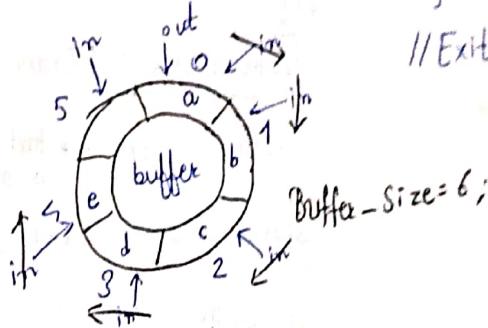
Consumer:

```
while(1) if counter==0
```

```
    while (in==out)
```

```
;  
nextConsumed = buffer[out],  
out = (out+1)%Buffer.size;
```

//Exit or Remainder Section



- Only (Buffer-Size-1) elements can be stored
- If we do (in == out) for producer as well, we won't be able to distinguish b/w empty queue and full queue
- Using another variable "counter" to keep track of the no. of elements in the queue, we can utilise the 1 wasted slot.

Producer:

```
reg1 = counter  
reg1 = reg1 + 1  
counter = reg1
```

Consumer:

```
reg2 = counter  
reg2 = reg2 - 1  
counter = reg2
```

- If both run sequentially, counter remains unchanged. (no issues)

- If both run concurrently (one possibility):

T₀: Producer reg₁ = counter; ⑤ reg₁

T₁: " reg₁ = reg₁ + 1; ⑥ reg₁

T₂: Consumer reg₂ = counter; ⑤ reg₂

T₃: " reg₂ = reg₂ - 1; ④ reg₂

T₄: Producer counter = reg₁; ⑥ counter

T₅: Consumer counter = reg₂; ④ counter

counter

5

6

4

(we run into issues)

This problem is called "Race Condition" (happens when multiple processes run concurrently sharing data)

- Part of code where shared resource is utilised, it is called "critical section".
- The idea is to not allow another process to enter its critical section, while one process is executing its critical section.

Semaphore

$s = 1$

`wait(s)` → waits if critical section is running

`signal(s)` → increments 's' if critical section isn't running

Semaphore

```
int s=1;
atomic {
    wait(s) {
        while (s <= 0)
            ; // No operation
        s--;
    }
    signal(s) {
        s++;
    }
}
```

* Mutual-exclusion implementation with semaphore

do {

`wait(mutex);` // Entry Section

critical section; → Part of code where shared data is accessed

`signal(mutex);` // Exit Section

remainder section;

? `while(1);`

// Code of Process: P;

- * Let $\{P_1, P_2, \dots, P_k\}$ processes want to access a shared data or resource, for the shared data there is a semaphore variable "mutex" which is initialized to 1.
- * `wait(s)` and `signal(s)` should be "atomic" functions, i.e. the function when called should be executed in its entirety or not at all, no in-between interruption.

* Disadvantage: if $s \leq 0$, the while loop will be constantly checking the condition, taking CPU resources (BUSY WAITING)

Solution:

```
typedef struct {
    int value;
    struct process * L;
}Semaphore;
```

`void wait(Semaphore s){`

`s.value --;`

`if (s.value < 0){`

 add this process to S.L;

 block();

(list of waiting processes)

`void signal(Semaphore s){`

`s.value ++;`

`if (S.value <= 0){`

 remove process P from S.L;

 Wakeup (P);

Synchronization Hardware

```

boolean Test-And-Set (boolean target) {
    boolean sv = target;
    target = true; // modifies the original
    if (target != sv) { // value of target not
        return sv; // a copy
    }
    // a copy
    // use address to modify
    // actual target
}

```

(NOT a function)
CPU instruction

```

do {
    while (Test-And-Set (lock)); // Entry Section
    critical section;
    lock = false; // Exit Section
} remainder section;

```

while (1);
 Mutual-exclusion Implementation with Test-And-Set
 condition — "Synchronization Hardware"
 To overcome race-around condition

Another technique satisfies three requirements:

Solution to Critical Section Problem should access critical section at a time

- i) Mutual Exclusion — only one process should access critical section at a time
 - ii) Progress — the processes not executing semantics
 - iii) Preemptive — the processes decide which decision made by process enters critical section next who're not in remainder section
- $p_1 \rightarrow$ Entry Section
 $p_2 \rightarrow$ Critical Section
 $p_3 \rightarrow$ Exit Section
Remainder Section
- iii) Preempted HardWare Waiting
 ↳ the processes should wait finitely
 If there are n processes, and 1 fails to get the CPU, it has to wait at most for $(n-1)$
 processes to execute.
- Synchronization Techniques have no control over which process gets the CPU first,
 a process may end up waiting indefinitely.

```

do{
    waiting[1] = true;
    key = true;
    while (waiting[1] && key) { Entry Section
        key = TestAndSet(lock);
        waiting[1] = false;
        // Critical Section
        j = (i+1)%m;
        while ((j != i) && !waiting[j]) { Critical Section
            j = (j+1)%m;
            if (j==1)
                lock=false;
            else
                waiting[j]=false;
                // Remaining Section
            ? while(1); // do nothing until lock is released
        } // End of Critical Section
        lock=true;
        waiting[1] = true;
    } // End of Entry Section
} // End of do loop

```

Common data structure used
 by P_0, P_1, \dots, P_{n-1} for lock
 is boolean waiting [m],
 boolean lock and return lock;

// Code for a process P_i